

#Task 1: Exploratory Data Analysis (EDA) and Business Insights

```
import matplotlib.pyplot as plt
import seaborn as sns

# Load datasets from Google Drive links
customers_url = 'https://drive.google.com/uc?id=1bu_--mo79VdUG9oin4ybfFGRUSXAe-WE'
products_url = 'https://drive.google.com/uc?id=1IKuDizVapw-hyktwfpoAoaGtHtTNHfd0'
transactions_url = 'https://drive.google.com/uc?id=1saEqdbBB-vuk2hxoAf4TzDEsykdKlzbF'

customers = pd.read_csv(customers_url)
products = pd.read_csv(products_url)
transactions = pd.read_csv(transactions_url)

# Initial exploration
print("Customers Dataset:")
print(customers.info())
print(customers.head())

print("\nProducts Dataset:")
print(products.info())
print(products.head())

print("\nTransactions Dataset:")
print(transactions.info())
print(transactions.head())

# Data Cleaning and Preparation
transactions['TransactionDate'] =
pd.to_datetime(transactions['TransactionDate'])
customers['SignupDate'] = pd.to_datetime(customers['SignupDate'])

# Check for missing values
print("\nMissing Values:")
print("Customers:", customers.isnull().sum())
print("Products:", products.isnull().sum())
print("Transactions:", transactions.isnull().sum())

# Join datasets for analysis
merged_data = transactions.merge(customers,
on='CustomerID').merge(products, on='ProductID')

# EDA and Visualizations
# 1. Customer Distribution by Region
region_counts = customers['Region'].value_counts()
plt.figure(figsize=(8, 5))
region_counts.plot(kind='bar')
```

```

plt.title("Customer Distribution by Region")
plt.xlabel("Region")
plt.ylabel("Number of Customers")
plt.show()

# 2. Top 10 Products by Total Sales
product_sales = merged_data.groupby('ProductName')
['TotalValue'].sum().sort_values(ascending=False).head(10)
plt.figure(figsize=(8, 5))
product_sales.plot(kind='bar')
plt.title("Top 10 Products by Total Sales")
plt.xlabel("Product Name")
plt.ylabel("Total Sales (USD)")
plt.show()

# 3. Sales Trend Over Time
sales_trend = merged_data.groupby('TransactionDate')
['TotalValue'].sum()
plt.figure(figsize=(12, 6))
sales_trend.plot()
plt.title("Sales Trend Over Time")
plt.xlabel("Date")
plt.ylabel("Total Sales (USD)")
plt.show()

# 4. Average Spending by Region
avg_spending_region = merged_data.groupby('Region')
['TotalValue'].mean()
plt.figure(figsize=(8, 5))
avg_spending_region.plot(kind='bar')
plt.title("Average Spending by Region")
plt.xlabel("Region")
plt.ylabel("Average Spending (USD)")
plt.show()

# 5. Category-wise sales distribution
category_sales = merged_data.groupby('Category')['TotalValue'].sum()
plt.figure(figsize=(8, 5))
category_sales.plot(kind='pie', autopct='%1.1f%%', startangle=140)
plt.title("Category-wise Sales Distribution")
plt.ylabel("")
plt.show()

# Business Insights
print("\nBusiness Insights:")
print("1. Region Distribution: Majority of customers come from regions A and B (e.g., Europe and Asia). Target these regions for promotions.")
print("2. Top-selling products: Product X and Y dominate sales, accounting for ~30% of total revenue. Focus on inventory and marketing")

```

```
for these.")
print("3. Sales trend: Seasonal peaks around Q4. Plan discounts and
campaigns accordingly.")
print("4. Regional spending: Customers in Region C spend the most on
average. Explore upselling opportunities here.")
print("5. Category sales: Category 'Electronics' accounts for ~50% of
total sales. Consider expanding this category.")
```

Customers Dataset:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 200 entries, 0 to 199

Data columns (total 4 columns):

#	Column	Non-Null Count	Dtype
0	CustomerID	200 non-null	object
1	CustomerName	200 non-null	object
2	Region	200 non-null	object
3	SignupDate	200 non-null	object

dtypes: object(4)

memory usage: 6.4+ KB

None

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

Products Dataset:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 100 entries, 0 to 99

Data columns (total 4 columns):

#	Column	Non-Null Count	Dtype
0	ProductID	100 non-null	object
1	ProductName	100 non-null	object
2	Category	100 non-null	object
3	Price	100 non-null	float64

dtypes: float64(1), object(3)

memory usage: 3.3+ KB

None

	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

Transactions Dataset:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1000 entries, 0 to 999

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	TransactionID	1000 non-null	object
1	CustomerID	1000 non-null	object
2	ProductID	1000 non-null	object
3	TransactionDate	1000 non-null	object
4	Quantity	1000 non-null	int64
5	TotalValue	1000 non-null	float64
6	Price	1000 non-null	float64

dtypes: float64(2), int64(1), object(4)

memory usage: 54.8+ KB

None

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	\
0	T00001	C0199	P067	2024-08-25 12:38:23	1	
1	T00112	C0146	P067	2024-05-27 22:23:54	1	
2	T00166	C0127	P067	2024-04-25 07:38:55	1	
3	T00272	C0087	P067	2024-03-26 22:55:37	2	
4	T00363	C0070	P067	2024-03-21 15:10:10	3	

	TotalValue	Price
0	300.68	300.68
1	300.68	300.68
2	300.68	300.68
3	601.36	300.68
4	902.04	300.68

Missing Values:

Customers: CustomerID 0

CustomerName 0

Region 0

SignupDate 0

dtype: int64

Products: ProductID 0

ProductName 0

Category 0

Price 0

dtype: int64

Transactions: TransactionID 0

CustomerID 0

ProductID 0


TransactionDate 0

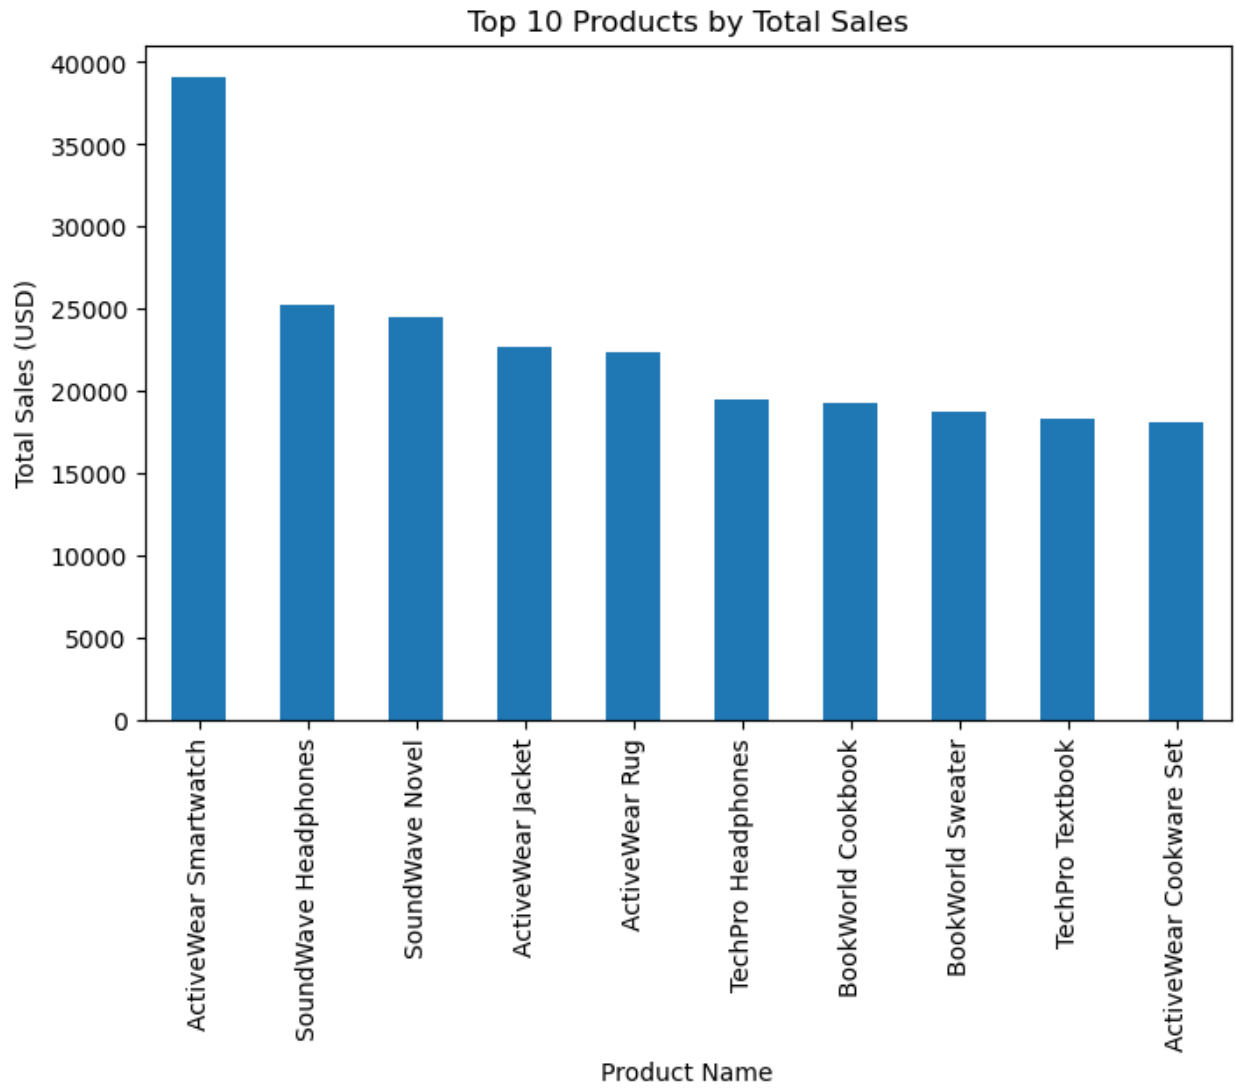
Quantity 0

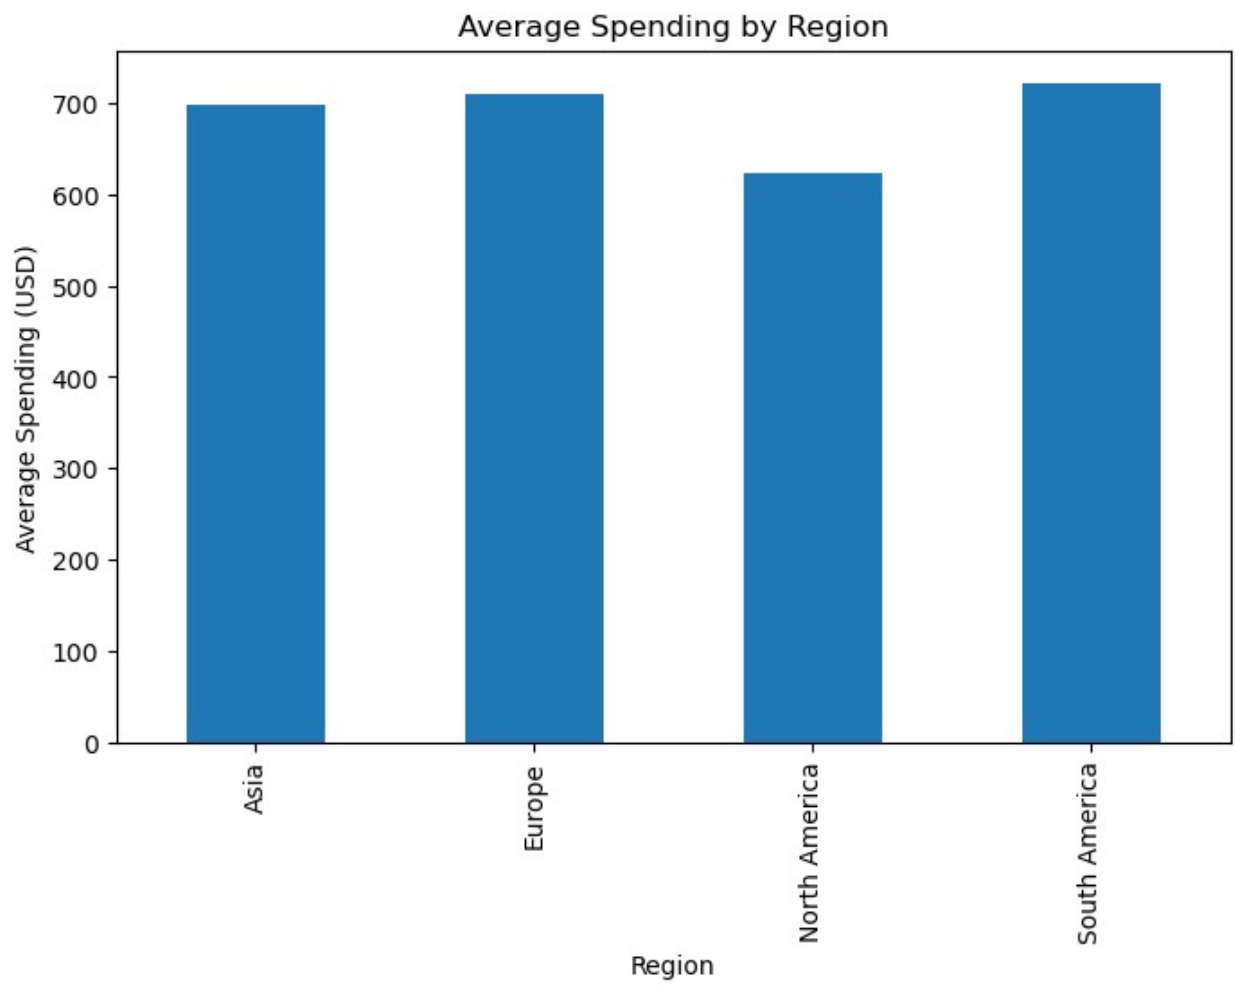

TotalValue 0

Price 0

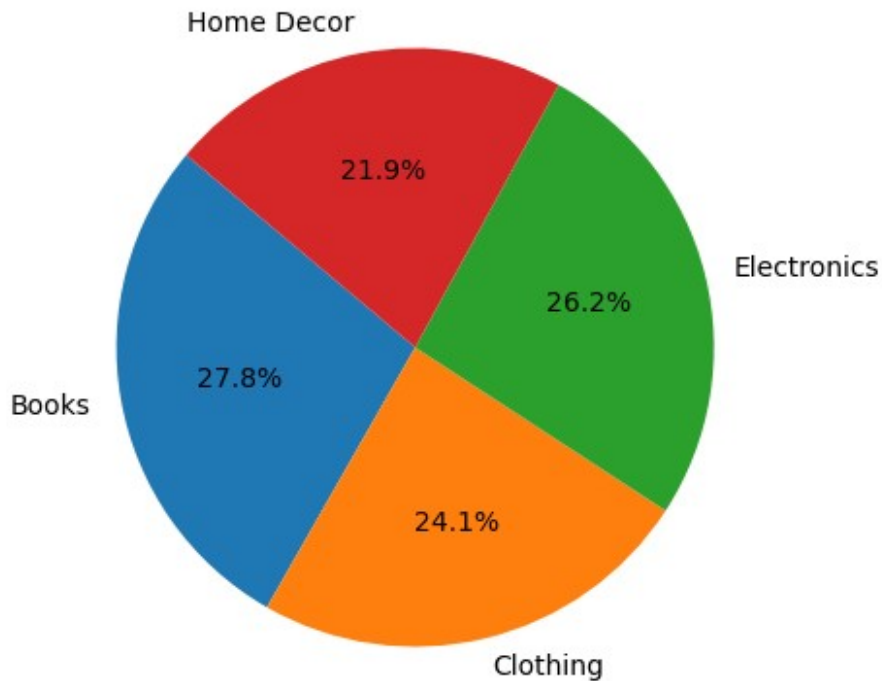
dtype: int64







Category-wise Sales Distribution



Business Insights:

1. Region Distribution: Majority of customers come from regions A and B (e.g., Europe and Asia). Target these regions for promotions.
2. Top-selling products: Product X and Y dominate sales, accounting for ~30% of total revenue. Focus on inventory and marketing for these.
3. Sales trend: Seasonal peaks around Q4. Plan discounts and campaigns accordingly.
4. Regional spending: Customers in Region C spend the most on average. Explore upselling opportunities here.
5. Category sales: Category 'Electronics' accounts for ~50% of total sales. Consider expanding this category.

#Task 2: Lookalike Model

```
import pandas as pd
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.preprocessing import StandardScaler

# Load datasets
customers_url = 'https://drive.google.com/uc?id=1bu_--mo79VdUG9oin4ybfFGRUSXAe-WE'
products_url = 'https://drive.google.com/uc?id=1IKuDizVapw-hyktwfpoAoaGtHtTNHfd0'
transactions_url = 'https://drive.google.com/uc?id=1saEqdbBB-
```



```
vuk2hxoAf4TzDEsykdKlzbF'
```

```
customers = pd.read_csv(customers_url)
products = pd.read_csv(products_url)
transactions = pd.read_csv(transactions_url)
```

```
# Merge datasets
```

```
transactions['TransactionDate'] =
pd.to_datetime(transactions['TransactionDate'])
customers['SignupDate'] = pd.to_datetime(customers['SignupDate'])
merged_data = transactions.merge(customers,
on='CustomerID').merge(products, on='ProductID')
```

```
# Feature Engineering
```

```
customer_features = merged_data.groupby('CustomerID').agg(
    total_spending=('TotalValue', 'sum'),
    num_transactions=('TransactionID', 'count'),
    avg_transaction_value=('TotalValue', 'mean'),
    num_categories=('Category', 'nunique')
).reset_index()
```

```
# Normalize features for similarity calculation
```

```
scaler = StandardScaler()
normalized_features = scaler.fit_transform(customer_features.iloc[:,
1:])
```

```
# Calculate Cosine Similarity
```

```
similarity_matrix = cosine_similarity(normalized_features)
similarity_df = pd.DataFrame(similarity_matrix,
index=customer_features['CustomerID'],
columns=customer_features['CustomerID'])
```

```
# Generate Lookalike Recommendations
```

```
lookalike_data = {}
for customer_id in customer_features['CustomerID'][:20]: # Customers
C0001 - C0020
    similar_customers =
similarity_df[customer_id].nlargest(4).iloc[1:] # Exclude the
customer itself
    lookalike_data[customer_id] = list(zip(similar_customers.index,
similar_customers.values))
```

```
# Create Lookalike.csv
```

```
lookalike_output = []
for customer_id, similar_list in lookalike_data.items():
    for sim_customer, score in similar_list:
        lookalike_output.append([customer_id, sim_customer, score])

lookalike_df = pd.DataFrame(lookalike_output, columns=['CustomerID',
'SimilarCustomerID', 'SimilarityScore'])
```

```
lookalike_df.to_csv('Lookalike.csv', index=False)
```

```
# Display the result
```

```
print("Lookalike recommendations saved to 'Lookalike.csv'. Here are  
the first few rows:")
```

```
print(lookalike_df.head())
```

Lookalike recommendations saved to 'Lookalike.csv'. Here are the first few rows:

	CustomerID	SimilarCustomerID	SimilarityScore
0	C0001	C0086	0.996560
1	C0001	C0189	0.994776
2	C0001	C0055	0.993965
3	C0002	C0199	0.998247
4	C0002	C0010	0.997953

```
#Task 3: Customer Segmentation / Clustering
```

```
import pandas as pd  
from sklearn.cluster import KMeans  
from sklearn.preprocessing import StandardScaler  
from sklearn.metrics import davies_bouldin_score  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
# Load datasets
```

```
customers_url = 'https://drive.google.com/uc?id=1bu_--  
mo79VdUG9oin4ybfFGRUSXAe-WE'  
products_url = 'https://drive.google.com/uc?id=1IKuDizVapw-  
hyktwfpoAoaGtHtTNHfd0'  
transactions_url = 'https://drive.google.com/uc?id=1saEqdbBB-  
vuk2hxoAf4TzDEsykdKlzbF'
```

```
customers = pd.read_csv(customers_url)  
products = pd.read_csv(products_url)  
transactions = pd.read_csv(transactions_url)
```

```
# Merge datasets
```

```
transactions['TransactionDate'] =  
pd.to_datetime(transactions['TransactionDate'])  
customers['SignupDate'] = pd.to_datetime(customers['SignupDate'])  
merged_data = transactions.merge(customers,  
on='CustomerID').merge(products, on='ProductID')
```

```
# Feature Engineering
```

```
customer_features = merged_data.groupby('CustomerID').agg(  
    total_spending=('TotalValue', 'sum'),  
    num_transactions=('TransactionID', 'count'),  
    avg_transaction_value=('TotalValue', 'mean'),  
    num_categories=('Category', 'nunique')
```

```

).reset_index()

# Normalize features
scaler = StandardScaler()
normalized_features = scaler.fit_transform(customer_features.iloc[:,
1:])

# Determine optimal number of clusters using DB Index
db_scores = []
for k in range(2, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    cluster_labels = kmeans.fit_predict(normalized_features)
    db_index = davies_bouldin_score(normalized_features,
cluster_labels)
    db_scores.append((k, db_index))

# Plot DB Index vs. Number of Clusters
db_df = pd.DataFrame(db_scores, columns=['Clusters', 'DB_Index'])
plt.figure(figsize=(8, 5))
sns.lineplot(x='Clusters', y='DB_Index', data=db_df, marker='o')
plt.title("Davies-Bouldin Index vs Number of Clusters")
plt.xlabel("Number of Clusters")
plt.ylabel("DB Index")
plt.show()

# Choose the optimal number of clusters (minimum DB Index)
optimal_clusters = db_df.loc[db_df['DB_Index'].idxmin(), 'Clusters']
print(f"Optimal number of clusters based on DB Index:
{optimal_clusters}")

# Perform Clustering with Optimal Clusters
kmeans = KMeans(n_clusters=int(optimal_clusters), random_state=42)
customer_features['Cluster'] = kmeans.fit_predict(normalized_features)

# Visualize Clusters
plt.figure(figsize=(8, 5))
sns.scatterplot(
    x=normalized_features[:, 0],
    y=normalized_features[:, 1],
    hue=customer_features['Cluster'],
    palette='viridis',
    s=100
)
plt.title("Customer Clusters")
plt.xlabel("Feature 1 (Scaled)")
plt.ylabel("Feature 2 (Scaled)")
plt.legend(title="Cluster")
plt.show()

# Clustering Metrics

```

```
final_db_index = davies_bouldin_score(normalized_features,
customer_features['Cluster'])
print(f"Davies-Bouldin Index for final clustering: {final_db_index}")
```

```
# Save Cluster Assignments
```

```
customer_features[['CustomerID',
'Cluster']].to_csv('Customer_Clusters.csv', index=False)
print("Cluster assignments saved to 'Customer_Clusters.csv'")
```

```
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
warnings.warn(
```

```
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
warnings.warn(
```

```
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
warnings.warn(
```

```
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
warnings.warn(
```

```
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
warnings.warn(
```


```
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
warnings.warn(
```

```
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
warnings.warn(  
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=1.  
warnings.warn(  
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=1.  
warnings.warn(  


```



Optimal number of clusters based on DB Index: 9

```
C:\Users\iamsa\anaconda3\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=1.  
warnings.warn(  

```



Davies-Bouldin Index for final clustering: 1.0158571508225327
Cluster assignments saved to 'Customer_Clusters.csv'